

YOLO

MobileNet

Vladislav Goncharenko
Fall 2024
MSAI

Recap

- Computer Vision tasks

- Classification
- Localisation
- Detection
- Semantic Segmentation
- Instance Segmentation

- Metrics

- IoU
- mAP

- Datasets

- PASCAL VOC
- ImageNet
- COCO
- Open Images




- R-CNN

- Fast R-CNN

- Faster R-CNN





Squeeze-and-Excitation Module



Outline



- Focal Loss
- Non-Maximum Suppression
- YOLO
 - v1
 - v2
 - V3
- Separable convolutions
- MobileNet

Focal Loss



[Video talk](#)

Non-Maximum Suppression



You Only Look Once (YOLO)



YOLOv1-v2

Original Joseph Redmon slides from conferences

- [YOLOv1](#)
- [YOLOv2](#)


Also [CVPR 2016 talk video](#)



YOLOv3

- Bounding Box Prediction
- Class Prediction
- Predictions Across Scales
- Feature Extractor

All articles published on [Joseph's website](#)



Practical notes


Write your own YOLO from scratch

<https://blog.paperspace.com/how-to-implement-a-yolo-object-detector-in-pytorch/>

Original YOLO written on Darknet - custom NN framework

Weights import to TF via

<https://github.com/thtrieu/darkflow>



Current state




Figure 5: Speed (ms) vs. accuracy (mAP) on COCO *test-dev*.

M2Det: A Single-Shot Object Detector based on Multi-Level Feature Pyramid Network

<https://arxiv.org/pdf/1811.04533.pdf>

Who is responsible for all these YOLOs

https://docs.google.com/spreadsheets/d/1BxI0SkyFGNvfzLLM_46xZ3i5crTst7jsUKKuZZDWrA0/edit#gid=0

Tools for detection

<https://github.com/roboflow/supervision>

Separable convs

MobileNet




Spatial Separable Convolution

$$\begin{bmatrix} 3 & 6 & 9 \\ 4 & 8 & 12 \\ 5 & 10 & 15 \end{bmatrix} = \begin{bmatrix} 3 \\ 4 \\ 5 \end{bmatrix} \times [1 \quad 2 \quad 3]$$

[Images credentials](#)

Spatial Separable Convolution



Simple Convolution




Spatial Separable Convolution



Normal Convolution




Depthwise Convolution



Don't change channels number

Pointwise Convolution



MobileNet v1






Table 1. MobileNet Body Architecture

Type / Stride	Filter Shape	Input Size
Conv / s2	$3 \times 3 \times 3 \times 32$	$224 \times 224 \times 3$
Conv dw / s1	$3 \times 3 \times 32$ dw	$112 \times 112 \times 32$
Conv / s1	$1 \times 1 \times 32 \times 64$	$112 \times 112 \times 32$
Conv dw / s2	$3 \times 3 \times 64$ dw	$112 \times 112 \times 64$
Conv / s1	$1 \times 1 \times 64 \times 128$	$56 \times 56 \times 64$
Conv dw / s1	$3 \times 3 \times 128$ dw	$56 \times 56 \times 128$
Conv / s1	$1 \times 1 \times 128 \times 128$	$56 \times 56 \times 128$
Conv dw / s2	$3 \times 3 \times 128$ dw	$56 \times 56 \times 128$
Conv / s1	$1 \times 1 \times 128 \times 256$	$28 \times 28 \times 128$
Conv dw / s1	$3 \times 3 \times 256$ dw	$28 \times 28 \times 256$
Conv / s1	$1 \times 1 \times 256 \times 256$	$28 \times 28 \times 256$
Conv dw / s2	$3 \times 3 \times 256$ dw	$28 \times 28 \times 256$
Conv / s1	$1 \times 1 \times 256 \times 512$	$14 \times 14 \times 256$
$5 \times$	Conv dw / s1 $3 \times 3 \times 512$ dw	$14 \times 14 \times 512$
	Conv / s1 $1 \times 1 \times 512 \times 512$	$14 \times 14 \times 512$
Conv dw / s2	$3 \times 3 \times 512$ dw	$14 \times 14 \times 512$
Conv / s1	$1 \times 1 \times 512 \times 1024$	$7 \times 7 \times 512$
Conv dw / s2	$3 \times 3 \times 1024$ dw	$7 \times 7 \times 1024$
Conv / s1	$1 \times 1 \times 1024 \times 1024$	$7 \times 7 \times 1024$
Avg Pool / s1	Pool 7×7	$7 \times 7 \times 1024$
FC / s1	1024×1000	$1 \times 1 \times 1024$
Softmax / s1	Classifier	$1 \times 1 \times 1000$


MobileNet v2




(a) Regular




(b) Separable



(c) Separable with linear bottleneck



(d) Bottleneck with expansion layer



Results

Architecture	Params	Top-1 accuracy	Top-5 accuracy
Xception	22.91M	0.790	0.945
VGG16	138.35M	0.715	0.901
MobileNetV1 (alpha=1, rho=1)	4.20M	0.709	0.899
MobileNetV1 (alpha=0.75, rho=0.85)	2.59M	0.672	0.873
MobileNetV1 (alpha=0.25, rho=0.57)	0.47M	0.415	0.663
MobileNetV2 (alpha=1.4, rho=1)	6.06M	0.750	0.925
MobileNetV2 (alpha=1, rho=1)	3.47M	0.718	0.910
MobileNetV2 (alpha=0.35, rho=0.43)	1.66M	0.455	0.704

MobileNet v3

Layer removal




Figure 5. Comparison of original last stage and efficient last stage. This more efficient last stage is able to drop three expensive layers at the end of the network at no loss of accuracy.

MobileNet v3

Nonlinearity

$$\text{swish } x = x \cdot \sigma(x)$$

$$\text{h-swish}[x] = x \frac{\text{ReLU6}(x + 3)}{6}$$





Figure 6. Sigmoid and swish nonlinearities and their “hard” counterparts.

Results

Network	Top-1	MAdds	Params	P-1	P-2	P-3
V3-Large 1.0	75.2	219	5.4M	51	61	44
V3-Large 0.75	73.3	155	4.0M	39	46	40
MnasNet-A1	75.2	315	3.9M	71	86	61
Proxyless[5]	74.6	320	4.0M	72	84	60
V2 1.0	72.0	300	3.4M	64	76	56
V3-Small 1.0	67.4	56	2.5M	15.8	19.4	14.4
V3-Small 0.75	65.4	44	2.0M	12.8	15.6	11.7
Mnas-small [43]	64.9	65.1	1.9M	20.3	24.2	17.2
V2 0.35	60.8	59.2	1.6M	16.6	19.6	13.9

Table 3. Floating point performance on the Pixel family of phones (P- n denotes a Pixel- n phone). All latencies are in ms and are measured using a single large core with a batch size of one. Top-1 accuracy is on ImageNet.

Revise

- Recap
- Focal Loss
- Non-Maximum Suppression
- YOLO
- Separable convolutions
- MobileNet

Next time

- Segmentation

See [Mask R-CNN ICCV17 talk](#) to prepare yourself

